

Assignment 1

Name: Barun Parua

Roll Number: 21CS10014

First of all, we take the excel input and convert it into a pandas dataframe for easier manipulation. We then apply normalisation and then split the data to appropriate subsets for training, validation and testing purposes.

As this is a binary classification problem, we define the class Ürgüp Sivrisi as 1 and the other class Çerçevelik as 0 for ease of notation. Note that this is just a convention.

```
In [ ]: # importing all the necessary libraries
# pandas for reading the csv file, numpy for mathematical operations, matplotlib
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [ ]: # extracting the dataset from the excel file, checking the shape of the dataset
df = pd.read_excel('../..../dataset/logistic-regression/Pumpkin_Seeds_Dataset.xlsx')
print("Shape of the dataset: ", df.shape)

# randomizing the dataset
df = df.sample(frac=1).reset_index(drop=True)

# splitting the dataset into x and y values
x_total = df.iloc[:, :-1].values
y_total = df.iloc[:, -1].values
print("Shape of x:", x_total.shape, "Shape of y:", y_total.shape)

y_total = np.where(y_total == 'Ürgüp Sivrisi', int(1), int(0))

print("First 5 rows of the dataset: \n", df.head())
print("\nFirst 5 values of y after converting to 0 and 1: \n", y_total[:5])
```

Shape of the dataset: (2500, 13)
 Shape of x: (2500, 12) Shape of y: (2500,)
 First 5 rows of the dataset:

	Area	Perimeter	Major_Axis_Length	Minor_Axis_Length	Convex_Area	\
0	70609	1011.299	393.8305	228.6390	71270	
1	93799	1285.858	542.4610	220.6602	94619	
2	77358	1147.854	488.9031	201.6832	78007	
3	65024	964.932	360.9694	230.4917	65576	
4	105209	1324.626	545.2486	246.9796	106238	

	Equiv_Diameter	Eccentricity	Solidity	Extent	Roundness	Aspect_Ration	\
0	299.8369	0.8142	0.9907	0.6798	0.8676	1.7225	
1	345.5844	0.9135	0.9913	0.6396	0.7129	2.4584	
2	313.8396	0.9109	0.9917	0.6538	0.7378	2.4241	
3	287.7345	0.7696	0.9916	0.7476	0.8776	1.5661	
4	366.0004	0.8915	0.9903	0.7217	0.7535	2.2077	

	Compactness	Class
0	0.7613	Çerçvelik
1	0.6371	Ürgüp Sivrisi
2	0.6419	Ürgüp Sivrisi
3	0.7971	Çerçvelik
4	0.6713	Ürgüp Sivrisi

First 5 values of y after converting to 0 and 1:
 [0 1 1 0 1]

```
In [ ]: # normalise the data using mean and standard deviation to avoid overflow and/or
for i in range(0, len(x_total[0])):
    mean = np.mean(x_total[:, i])
    std = np.std(x_total[:, i])
    x_total[:, i] = (x_total[:, i] - mean) / std
```

```
In [ ]: # split the data into train, test and validation set
# train set: 50%
# validation set: 30%
# test set: 20%
train_size = int(0.5 * len(x_total))
validation_size = int(0.3 * len(x_total))
test_size = int(0.2 * len(x_total))

# splitting the data into train, validation and test set
x_train = x_total[:train_size]
x_validation = x_total[train_size:train_size+validation_size]
x_test = x_total[train_size+validation_size:]

y_train = y_total[:train_size]
y_validation = y_total[train_size:train_size+validation_size]
y_test = y_total[train_size+validation_size:]

# checking the shape of the data
print("Shape of x_train: ", x_train.shape)
print("Shape of y_train: ", y_train.shape)

print("Shape of x_test: ", x_test.shape)
print("Shape of y_test: ", y_test.shape)
```

```

Shape of x_train: (1250, 12)
Shape of y_train: (1250,)
Shape of x_test: (500, 12)
Shape of y_test: (500,)

```

Now, as this is a binary classification problem, we will use the sigmoid function as our hypothesis function. Then on the basis of the hypothesis function, we predict the classes of the data points. Then we defined the appropriate cost function and then used gradient descent method to minimize the costs.

```

In [ ]: # gradient descent method for Logistic regression
        # necessary functions for gradient descent method

        # sigmoid and predict functions
        def sigmoid(z):
            return 1 / (1 + np.exp(-z))

        def predict(x, theta):
            return sigmoid(np.dot(x, theta))

        # cost function and gradient descent function
        def cost(x, y, theta):
            return -np.mean(y * np.log(predict(x, theta)) + (1 - y) * np.log(1 - predict(x, theta)))

        def gradient_descent(x, y, theta, learning_rate):
            return theta - learning_rate * np.dot(x.T, predict(x, theta) - y) / len(y)

```

Finally, we do the Gradient Descent and find the optimal values of the parameters theta using different learning rates. It can be seen that the number of iterations is less for higher learning rates. This is because the step size is larger for higher learning rates and hence the convergence is faster. Therefore I have set it such that number of iterations for a learning rate is $10 * (1/\text{learning_rate_value})$.

It can be seen that the cost value attains a minimum value in each case as we have set the iterations properly. Hence, again the parameters are used to predict the values of y for the test data.

Now while the output y can be any value between 0 and 1, we need to convert it to 0 or 1. This is done by setting a threshold value. If the value of y is greater than the threshold value, then it is set to 1, else it is set to 0. The threshold value is set to 0.5 for this case.

Finally, the accuracy, precision, recall and are calculated for each of the learning rates. The results for all the learning rates are shown below.

```

In [ ]: # applying gradient descent on the train set for appropriate number of iteration
        theta_grad = np.random.rand(x_train.shape[1])
        costs = []
        learning_rate = 0.1
        for i in range(100):
            costs.append(cost(x_train, y_train, theta_grad))
            theta_grad = gradient_descent(x_train, y_train, theta_grad, learning_rate)
        print("Analysis for Learning Rate =", learning_rate, "with", i+1, "iterations: ")
        print("Theta: ")
        print(theta_grad)

```

```

# apply the obtained theta on the test set and calculate the accuracy, precision
y_pred = np.dot(x_test, theta_grad)

# calculating true positives, true negatives, false positives and false negative
tp, tn, fp, fn, = 0, 0, 0, 0

# apply the threshold of 0.5 and check the values of y_pred and y_test
for i in range(len(y_pred)):
    if y_pred[i] >= 0.5:
        y_pred[i] = 1
    else:
        y_pred[i] = 0
    if y_pred[i] == 1 and y_test[i] == 1:
        tp += 1
    elif y_pred[i] == 0 and y_test[i] == 0:
        tn += 1
    elif y_pred[i] == 1 and y_test[i] == 0:
        fp += 1
    else:
        fn += 1

# calculating accuracy, precision and recall based on the obtained values
print("Accuracy: ", (tp+tn)/(tp+tn+fp+fn))
print("Precision: ", tp/(tp+fp))
print("Recall: ", tp/(tp+fn))

# plot the cost using matplotlib to get an idea of how the cost function varies
plt.plot(costs)
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.title('Cost vs Iterations (Learning Rate = 0.1)')
plt.show()

```

Analysis for Learning Rate = 0.1 with 100 iterations:

Theta:

```


[-0.13986518 -0.03279262  0.40759002 -0.06040112 -0.13519599  0.33809902
  0.72816883  0.22244932 -0.00282432 -0.13680537  1.32062643 -0.39662375]

```

Accuracy: 0.872

Precision: 0.9319371727748691

Recall: 0.777292576419214



```
In [ ]: # applying gradient descent on the train set for appropriate number of iteration
theta_grad = np.random.rand(x_train.shape[1])
costs = []
learning_rate = 0.01
for i in range(1000):
    costs.append(cost(x_train, y_train, theta_grad))
    theta_grad = gradient_descent(x_train, y_train, theta_grad, learning_rate)
print("Analysis for Learning Rate =", learning_rate, "with", i+1, "iterations: ")
print("Theta: ")
print(theta_grad)

# apply the obtained theta on the test set and calculate the accuracy, precision
y_pred = np.dot(x_test, theta_grad)

# calculating true positives, true negatives, false positives and false negative
tp, tn, fp, fn, = 0, 0, 0, 0

# apply the threshold of 0.5 and check the values of y_pred and y_test
for i in range(len(y_pred)):
    if y_pred[i] >= 0.5:
        y_pred[i] = 1
    else:
        y_pred[i] = 0
    if y_pred[i] == 1 and y_test[i] == 1:
        tp += 1
    elif y_pred[i] == 0 and y_test[i] == 0:
        tn += 1
    elif y_pred[i] == 1 and y_test[i] == 0:
        fp += 1
    else:
        fn += 1
```

```
# calculating accuracy, precision and recall based on the obtained values
print("Accuracy: ", (tp+tn)/(tp+tn+fp+fn))
print("Precision: ", tp/(tp+fp))
print("Recall: ", tp/(tp+fn))

# plot the cost using matplotlib to get an idea of how the cost function varies
plt.plot(costs)
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.title('Cost vs Iterations (Learning Rate = 0.01)')
plt.show()
```

Analysis for Learning Rate = 0.01 with 1000 iterations:


Theta:

```
[ 0.56469085  0.31537608  0.32384319 -0.1166017  -0.31744793 -0.40916841
  0.80574735  0.39710624  0.14571165 -0.46417713  0.89289562 -0.31849838]
```

Accuracy: 0.868

Precision: 0.9267015706806283

Recall: 0.7729257641921398



```
In [ ]: # applying gradient descent on the train set for appropriate number of iteration
theta_grad = np.random.rand(x_train.shape[1])
costs = []
learning_rate = 0.001
for i in range(10000):
    costs.append(cost(x_train, y_train, theta_grad))
    theta_grad = gradient_descent(x_train, y_train, theta_grad, learning_rate)
print("Analysis for Learning Rate =", learning_rate, "with", i+1, "iterations: ")
print("Theta: ")
print(theta_grad)

# apply the obtained theta on the test set and calculate the accuracy, precision
y_pred = np.dot(x_test, theta_grad)
```

```

# calculating true positives, true negatives, false positives and false negative
tp, tn, fp, fn, = 0, 0, 0, 0

# apply the threshold of 0.5 and check the values of y_pred and y_test
for i in range(len(y_pred)):
    if y_pred[i] >= 0.5:
        y_pred[i] = 1
    else:
        y_pred[i] = 0
    if y_pred[i] == 1 and y_test[i] == 1:
        tp += 1
    elif y_pred[i] == 0 and y_test[i] == 0:
        tn += 1
    elif y_pred[i] == 1 and y_test[i] == 0:
        fp += 1
    else:
        fn += 1

# calculating accuracy, precision and recall based on the obtained values
print("Accuracy: ", (tp+tn)/(tp+tn+fp+fn))
print("Precision: ", tp/(tp+fp))
print("Recall: ", tp/(tp+fn))

# plot the cost using matplotlib to get an idea of how the cost function varies
plt.plot(costs)
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.title('Cost vs Iterations (Learning Rate = 0.001)')
plt.show()

```

Analysis for Learning Rate = 0.001 with 10000 iterations:

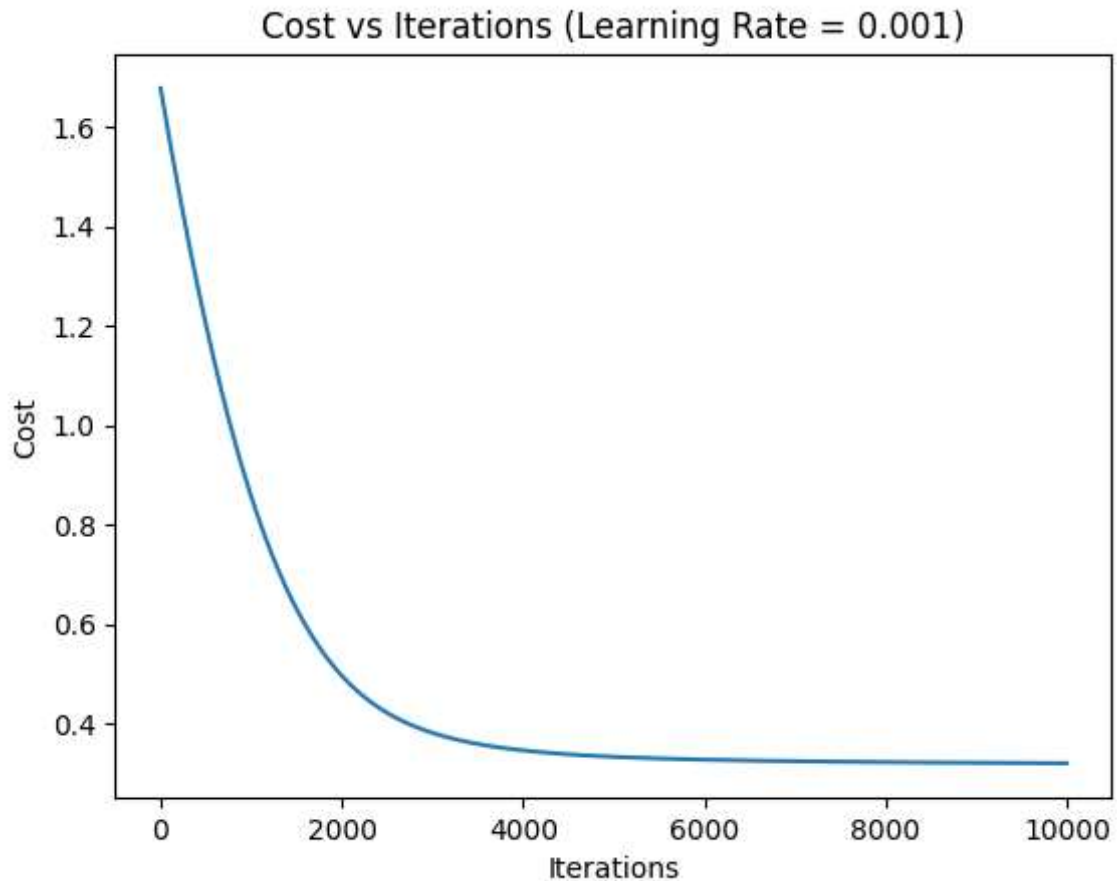
Theta:

```
[ 0.08306153  0.19597722 -0.02652305 -0.29636562  0.10977027  0.1418317
 1.33655143  0.45496433  0.21116648 -0.61721659  0.81056922  0.03249813]
```

Accuracy: 0.874

Precision: 0.9191919191919192

Recall: 0.7947598253275109



```
In [ ]: # applying gradient descent on the train set for appropriate number of iteration
theta_grad = np.random.rand(x_train.shape[1])
costs = []
learning_rate = 0.0001
for i in range(100000):
    costs.append(cost(x_train, y_train, theta_grad))
    theta_grad = gradient_descent(x_train, y_train, theta_grad, learning_rate)
print("Analysis for Learning Rate =", learning_rate, "with", i+1, "iterations: ")
print("Theta: ")
print(theta_grad)

# apply the obtained theta on the test set and calculate the accuracy, precision
y_pred = np.dot(x_test, theta_grad)

# calculating true positives, true negatives, false positives and false negative
tp, tn, fp, fn, = 0, 0, 0, 0

# apply the threshold of 0.5 and check the values of y_pred and y_test
for i in range(len(y_pred)):
    if y_pred[i] >= 0.5:
        y_pred[i] = 1
    else:
        y_pred[i] = 0
    if y_pred[i] == 1 and y_test[i] == 1:
        tp += 1
    elif y_pred[i] == 0 and y_test[i] == 0:
        tn += 1
    elif y_pred[i] == 1 and y_test[i] == 0:
        fp += 1
    else:
        fn += 1
```



```
# calculating accuracy, precision and recall based on the obtained values
print("Accuracy: ", (tp+tn)/(tp+tn+fp+fn))
print("Precision: ", tp/(tp+fp))
print("Recall: ", tp/(tp+fn))

# plot the cost using matplotlib to get an idea of how the cost function varies
plt.plot(costs)
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.title('Cost vs Iterations (Learning Rate = 0.0001)')
plt.show()
```

Analysis for Learning Rate = 0.0001 with 100000 iterations:

Theta:

```
[-0.3008083  0.63760531  0.82657353 -0.18495087 -0.22237932 -0.21110172
 1.22886002  0.57747896 -0.0077319  0.36594054  1.01689921 -0.23113656]
```

Accuracy: 0.874

Precision: 0.915

Recall: 0.7991266375545851

