

# CS60050

# Machine Learning

## Decision Trees

Somak Aditya

Assistant Professor  
Department of CSE  
August 16, 2023

Sudeshna Sarkar

Professor  
CoE AI, Department of CSE



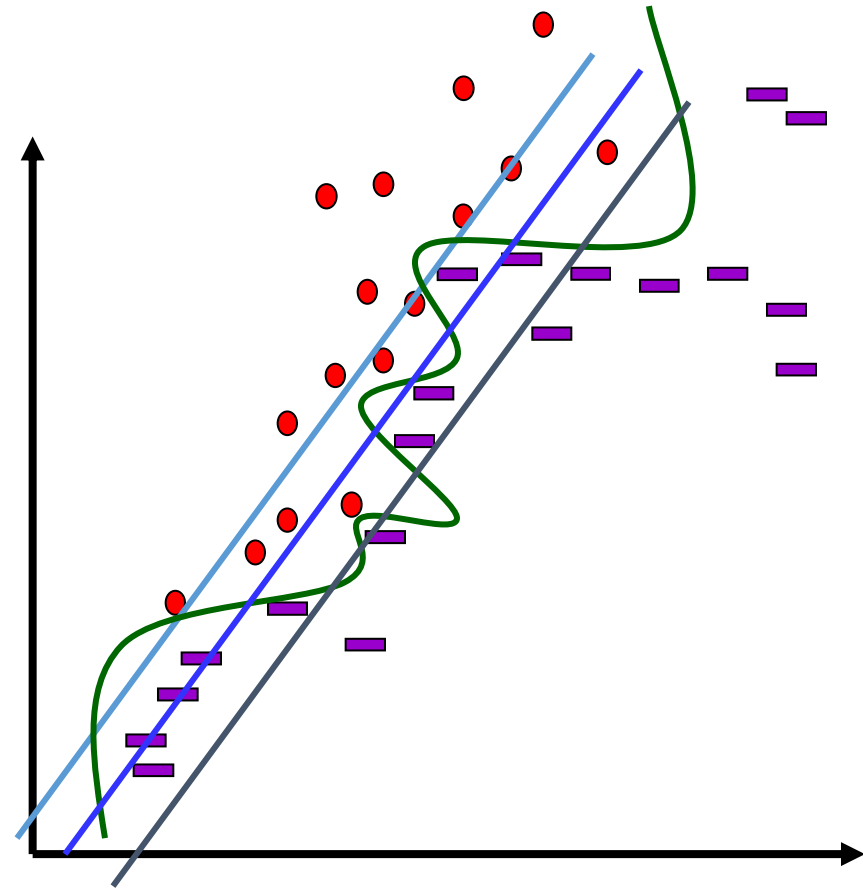
# Recap

- Learning problem (classification)  
Find a function that  
best separates the data
- What function?

Linear:

$$Y = \text{sign}(\theta^T X)$$

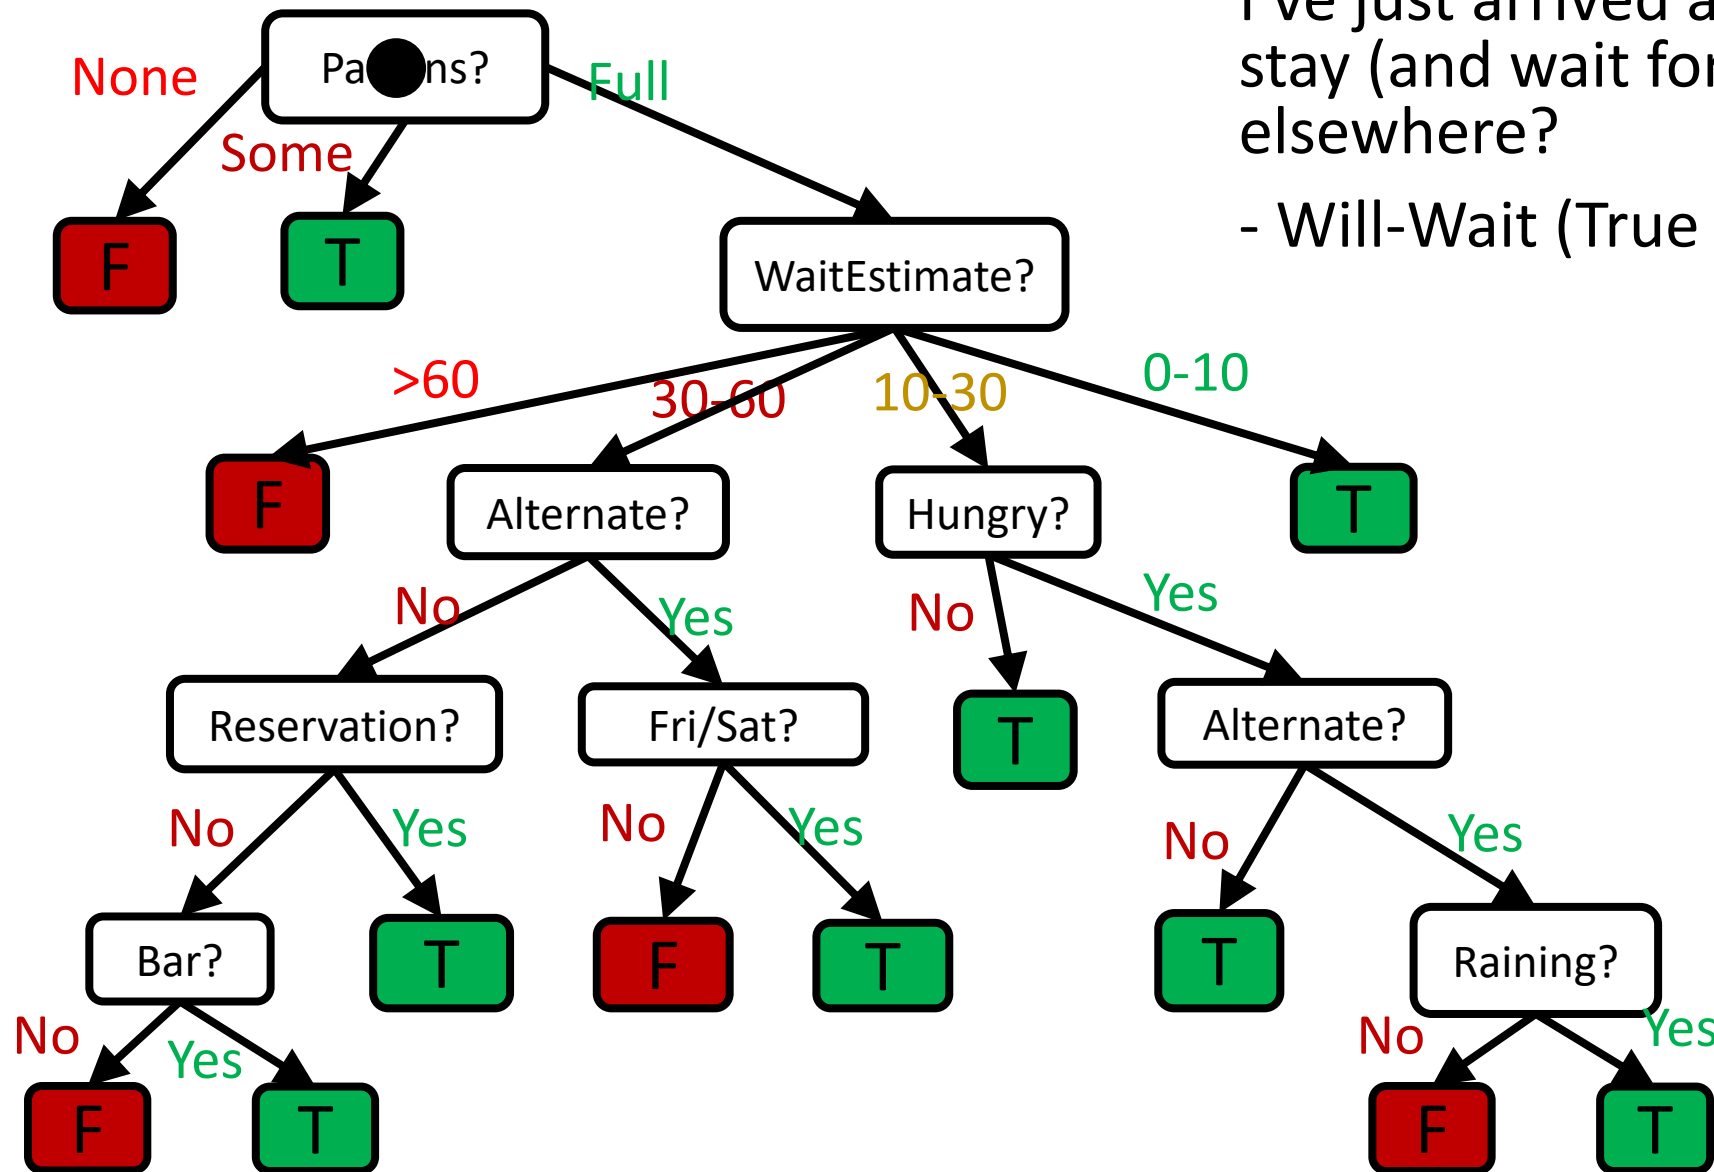
- Limitations of Linear Functions



# Decision Trees

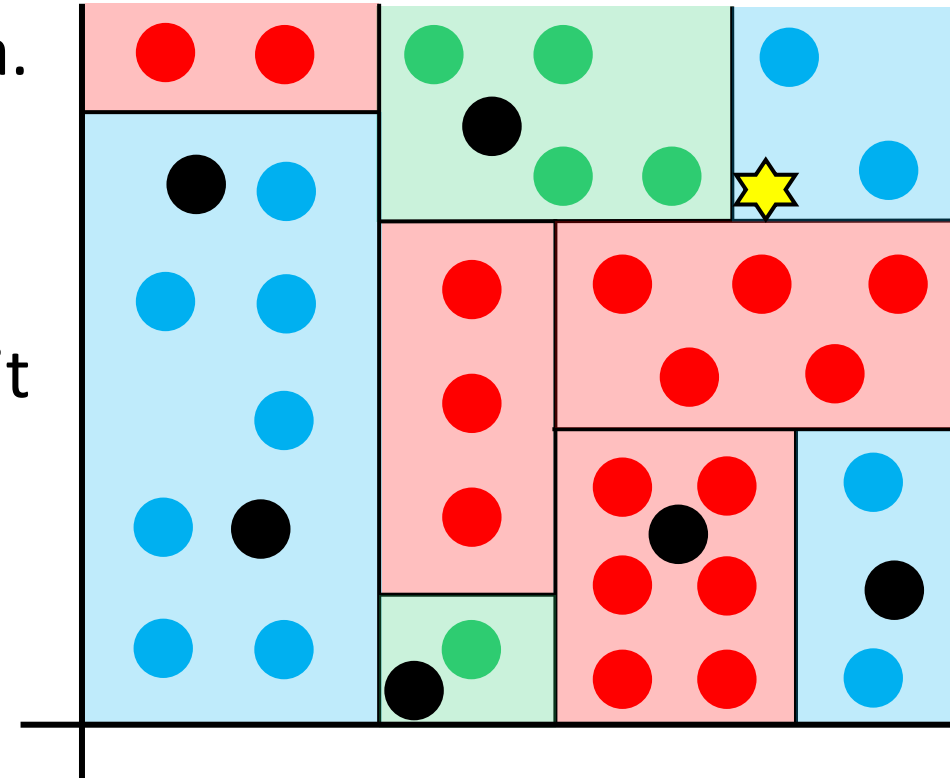
I've just arrived at a restaurant: should I stay (and wait for a table) or go elsewhere?

- Will-Wait (True or False?)



# Decision Trees for Classification

- Find "many" lines that best "separates" the data.
- Repeatedly **partition** feature space  $\mathbb{R}^d$
- Assign a label to each partition.
- For test data point, easy to find which partition it lies

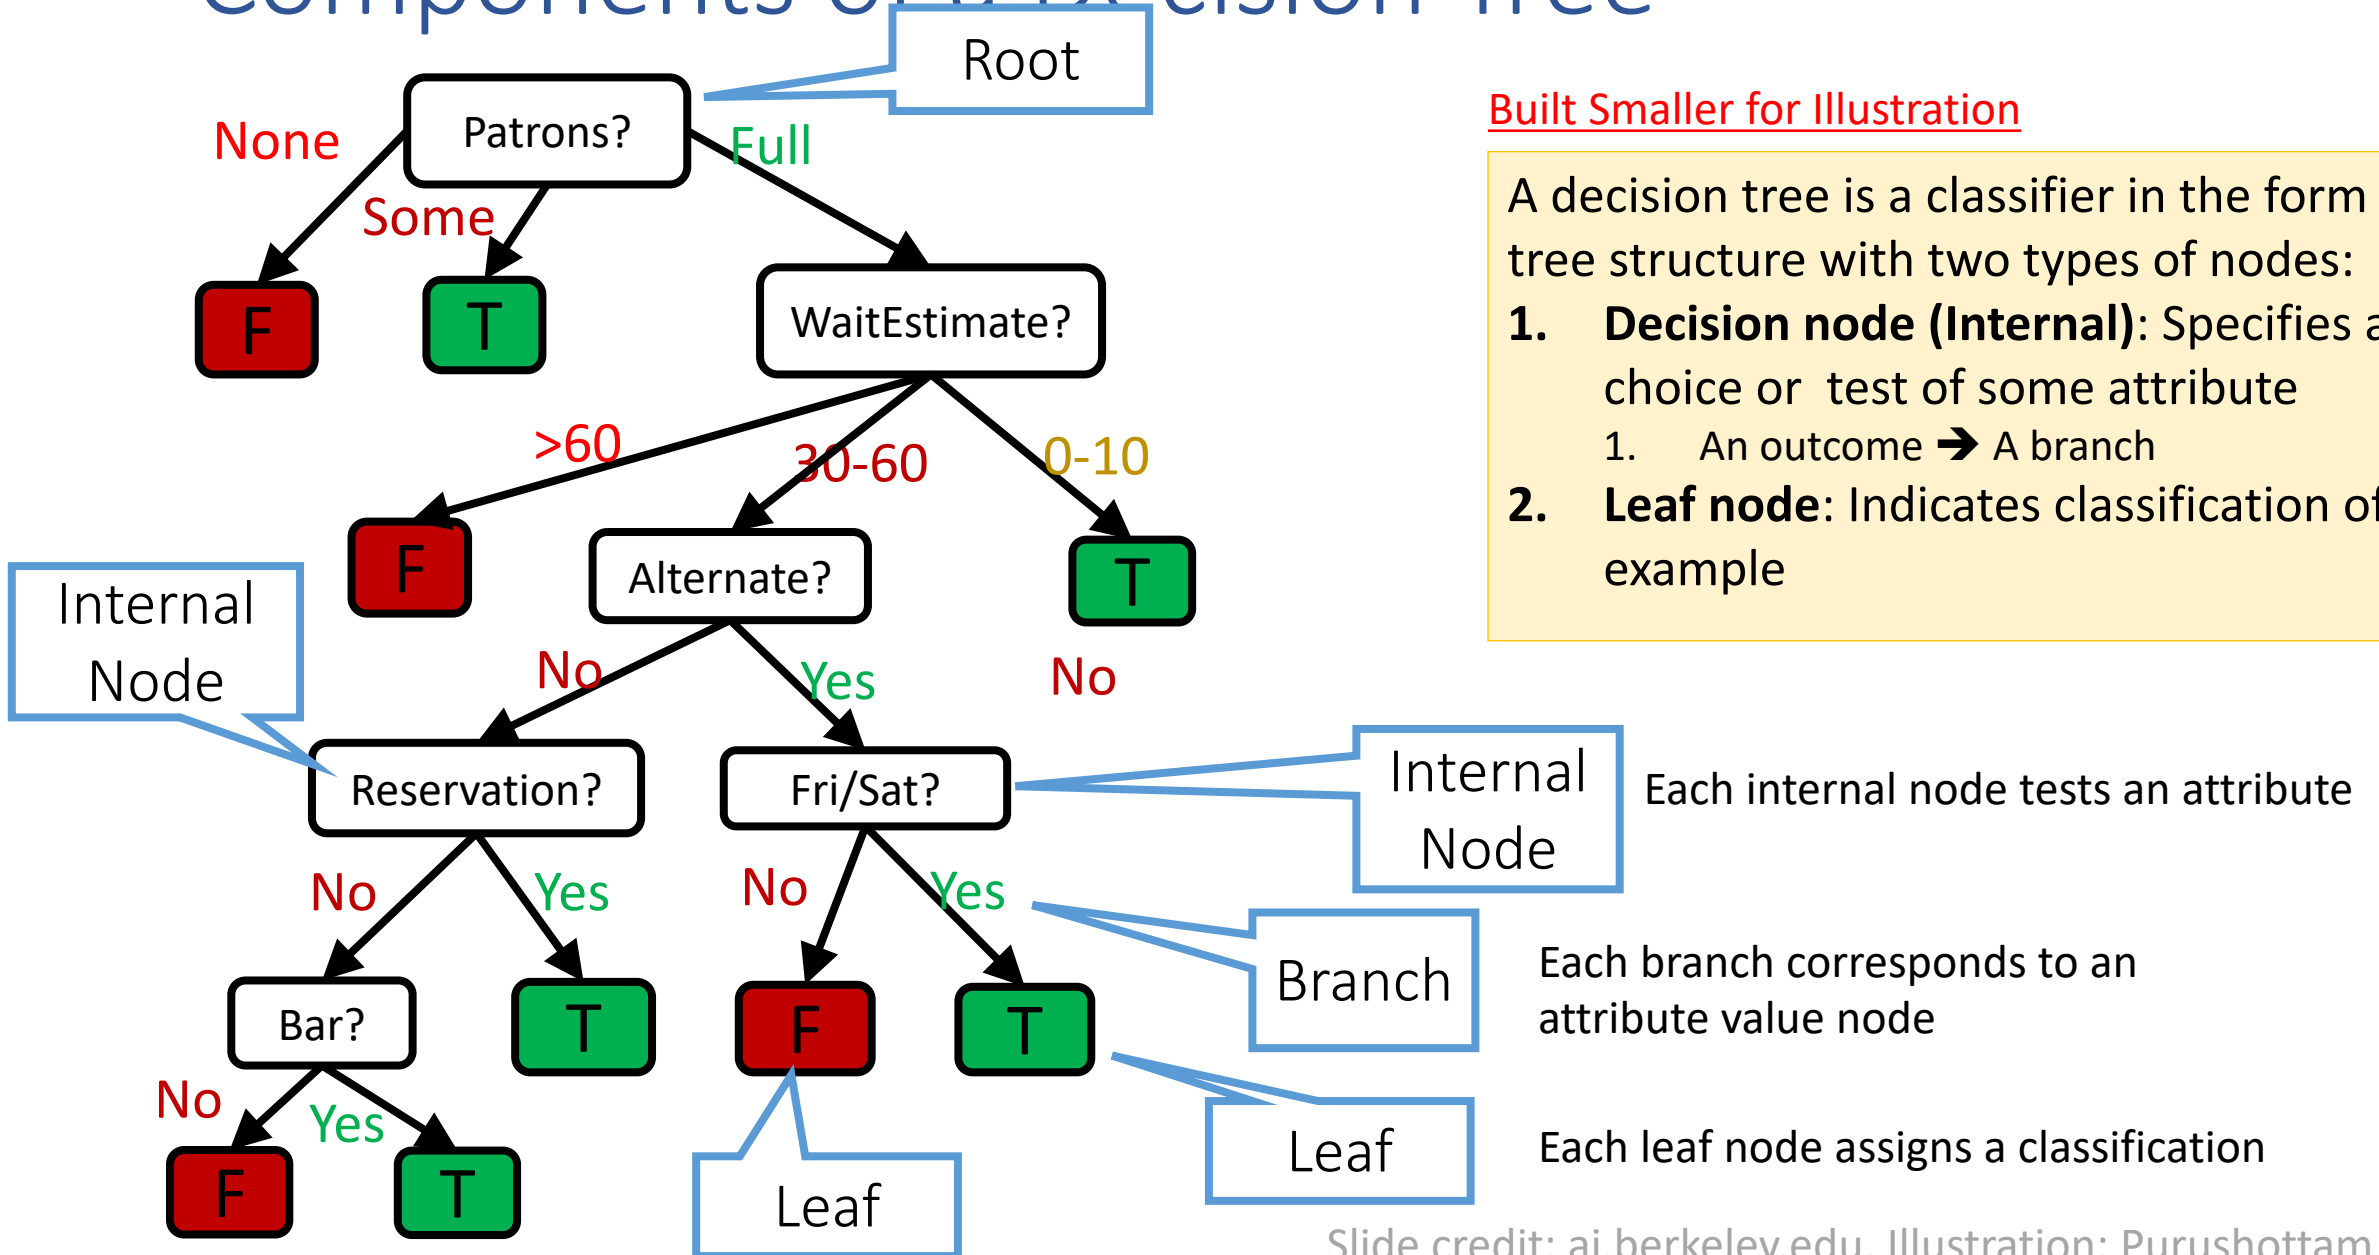


# Components of a Decision Tree

Built Smaller for Illustration

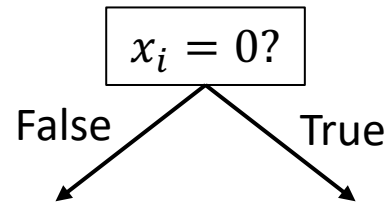
A decision tree is a classifier in the form of a tree structure with two types of nodes:

1. **Decision node (Internal):** Specifies a choice or test of some attribute
  1. An outcome → A branch
2. **Leaf node:** Indicates classification of an example

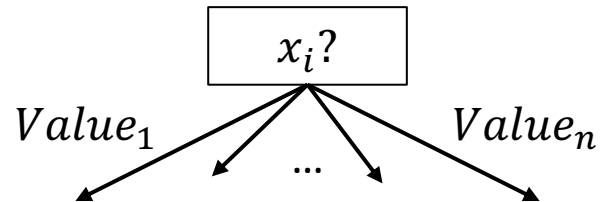


## Type of Internal Tests

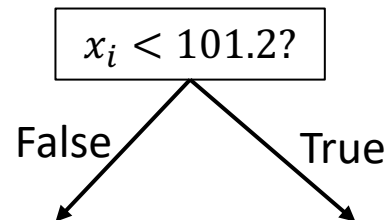
Binary Feature



Categorical Feature



Numeric Feature



## Types of Leaves

Classification

$$y = POS$$

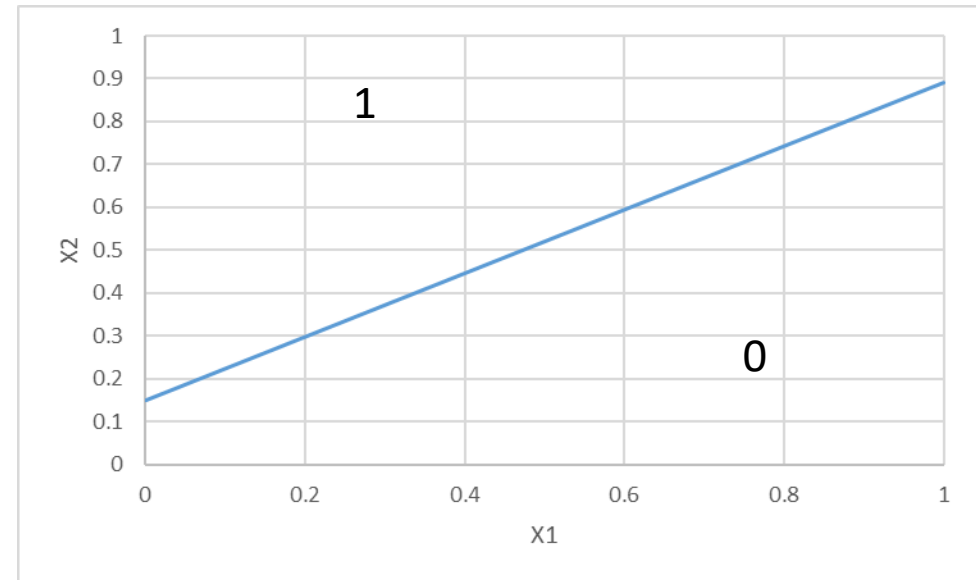
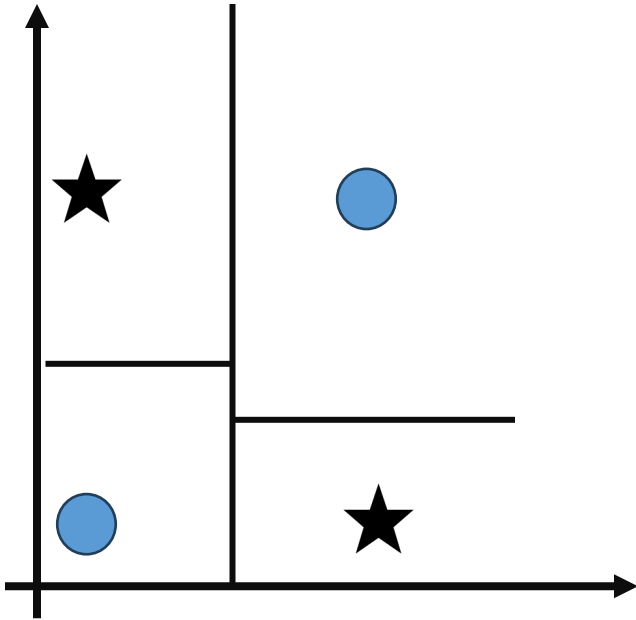
Regression

$$y = 5.2$$

Probability Estimate

$$\begin{aligned} p(y = 0) &= .3 \\ p(y = 1) &= .3 \\ p(y = 2) &= .4 \end{aligned}$$

# Decision Trees vs Linear Models



Adding nodes (structure) allows more powerful representation

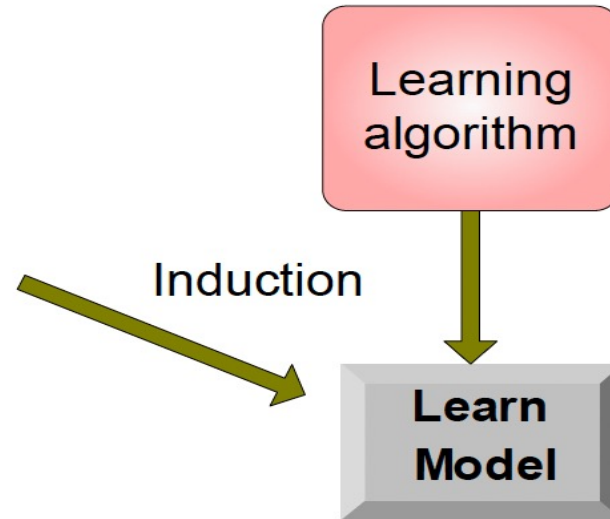
# Classification in a Nutshell

<i>Tid</i>	<i>Attrib1</i>	<i>Attrib2</i>	<i>Attrib3</i>	<i>Class</i>
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

<i>Tid</i>	<i>Attrib1</i>	<i>Attrib2</i>	<i>Attrib3</i>	<i>Class</i>
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set





# Learning Decision Tree from Data

<i>Tid</i>	<i>Refund</i>	<i>Marital Status</i>	<i>Taxable Income</i>	<i>Cheat</i>
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

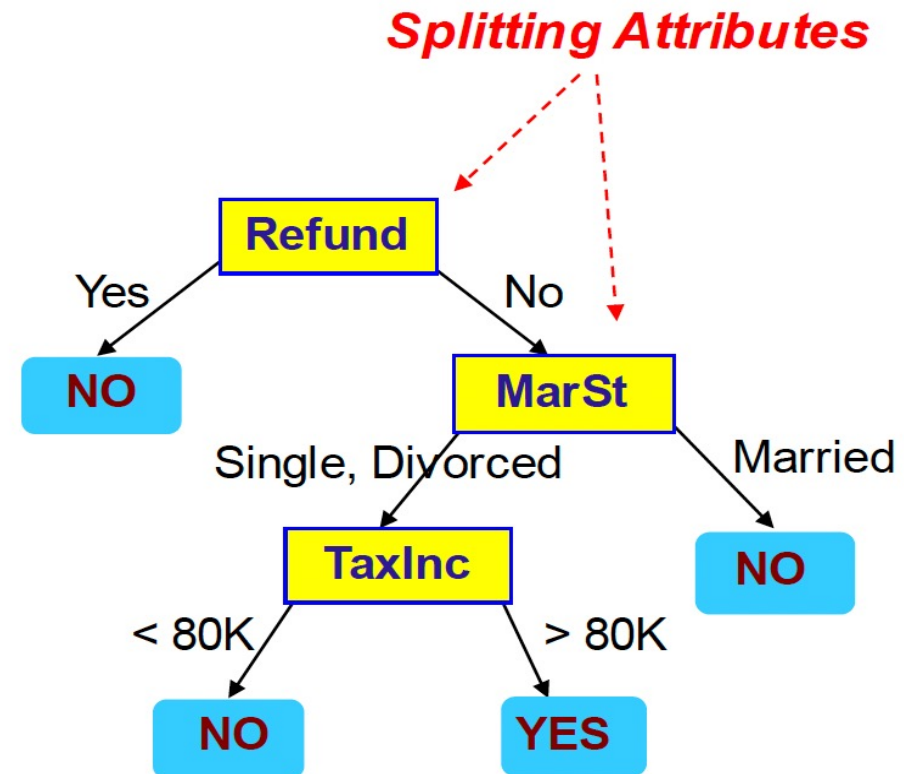
categorical

categorical

continuous

class

Training Data



Model: Decision Tree

# Issues

- Given some training examples, what decision tree should be generated?
- One proposal: learn the smallest tree that is *fits the data*
- Possible method:
  - Exhaustive search over the space of decision trees
  - This is NP-hard.
- *Efficient algorithms available to learn a **reasonably accurate** (potentially suboptimal) decision tree in **reasonable time***
  - ***Employs greedy strategy***
  - *Locally optimal choices about which attribute to use next to partition the data*

# Building a Decision Tree (GREEDY)

Function `BuildTree(dataset, attributes)` returns a DT

# *dataset* : dataset at current node, *attributes*: current set of attributes

If *attributes* is empty OR

all labels in *dataset* are the same:

# Leaf node

class = most common class in *dataset*

else

# Internal node

*att*  $\leftarrow$  CHOOSE-BEST-ATTRIBUTE(*dataset*, *attributes*)

*tree*  $\leftarrow$  A new DT with root test *att*

LeftNode = BuildTree(*dataset*(*att* = 1), *attributes* - {*att*})

RightNode = BuildTree(*dataset*(*att* = 0), *attributes* - {*att*})

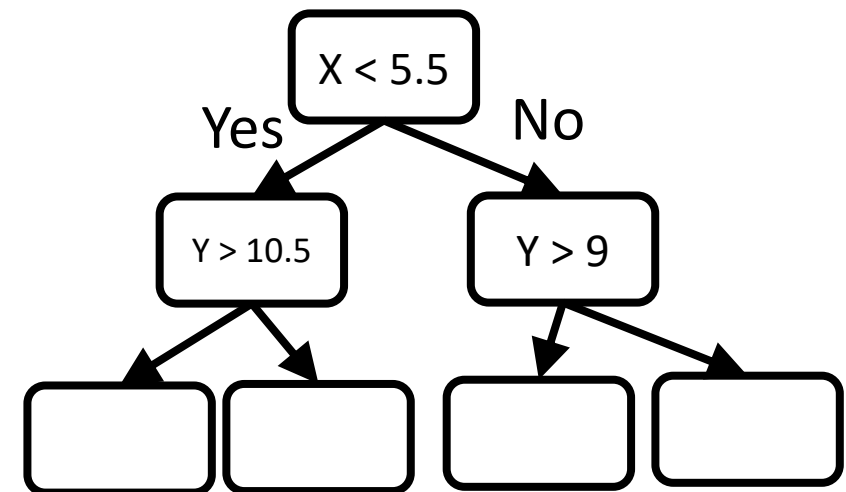
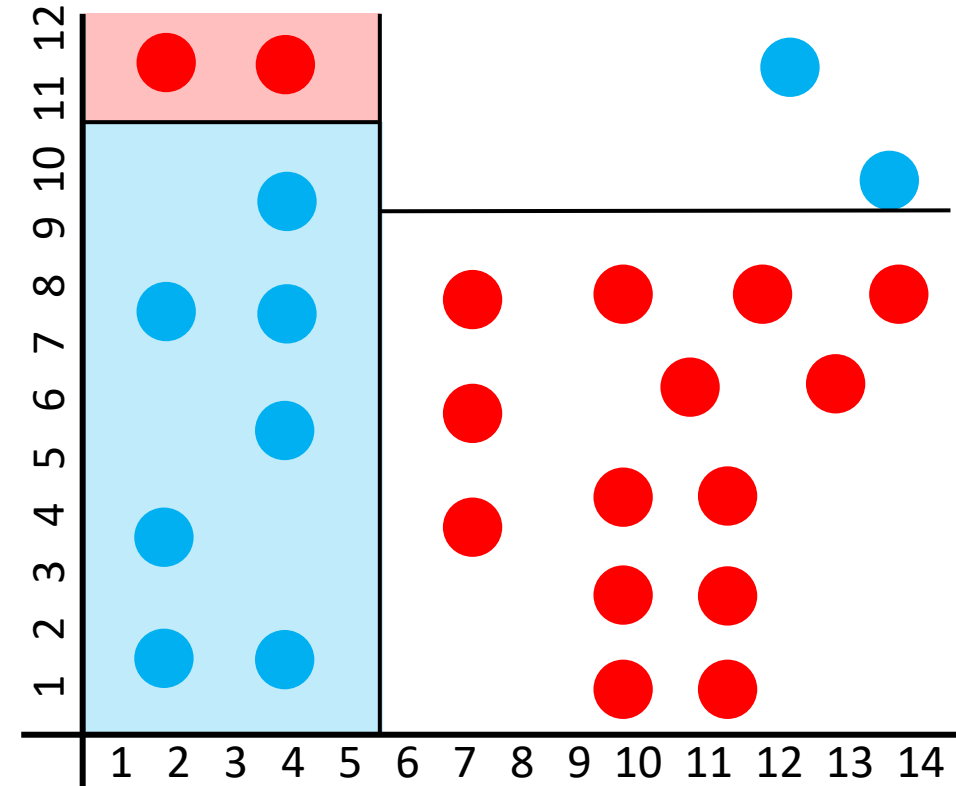
add branch to *tree* with value 1, subtree LeftNode

add branch to *tree* with value 0, subtree RightNode

# generalize for multiple values

return *tree*

Binary  
Test



# Decision Tree (Learn & Predict)

Function BuildTree(dataset,attributes) returns a DT

# *dataset* : dataset at current node, *attributes*: current set of attributes

If *attributes* is empty OR

all labels in *dataset* are the same:

# Leaf node

class = most common class in *dataset*

else

# Internal node

*att*  $\leftarrow$  CHOOSE-BEST-ATTRIBUTE(*dataset*, *attributes*)

*tree*  $\leftarrow$  A new DT with root test *att*

For each value  $v_j$  of attribute *att*:

ChildNode<sub>*j*</sub> = BuildTree(D(*att*=  $v_j$ ), attributes—{*att*})

add branch to *tree* with value  $v_j$ , subtree ChildNode<sub>*j*</sub>

return *tree*

## Prediction

def  $f(x')$ :

Let *current node* = root

while(true):

- if *current node* is internal (non-leaf):
  - Let *a* = attribute associated with current node
  - Go down branch labeled with value  $x'_a$
- if *current node* is a leaf:
  - return label *y* stored at that leaf

# Decision Tree (Choices)

Function `BuildTree(dataset, attributes)` returns a DT

# *dataset* : dataset at current node, *attributes*: current set of attributes

If *attributes* is empty OR

all labels in *dataset* are the same:

# Leaf node

class = most common class in *dataset*

else

# Internal node

*att*  $\leftarrow$  CHOOSE-BEST-ATTRIBUTE(*dataset*, *attributes*)

*tree*  $\leftarrow$  A new DT with root test *att*

For each value  $v_j$  of attribute *att*:

$\text{ChildNode}_j = \text{BuildTree}(D(\text{att} = v_j), \text{attributes} - \{\text{att}\})$

    add branch to *tree* with value  $v_j$ , subtree  $\text{ChildNode}_j$

return *tree*

## Choices

### 1. When to stop

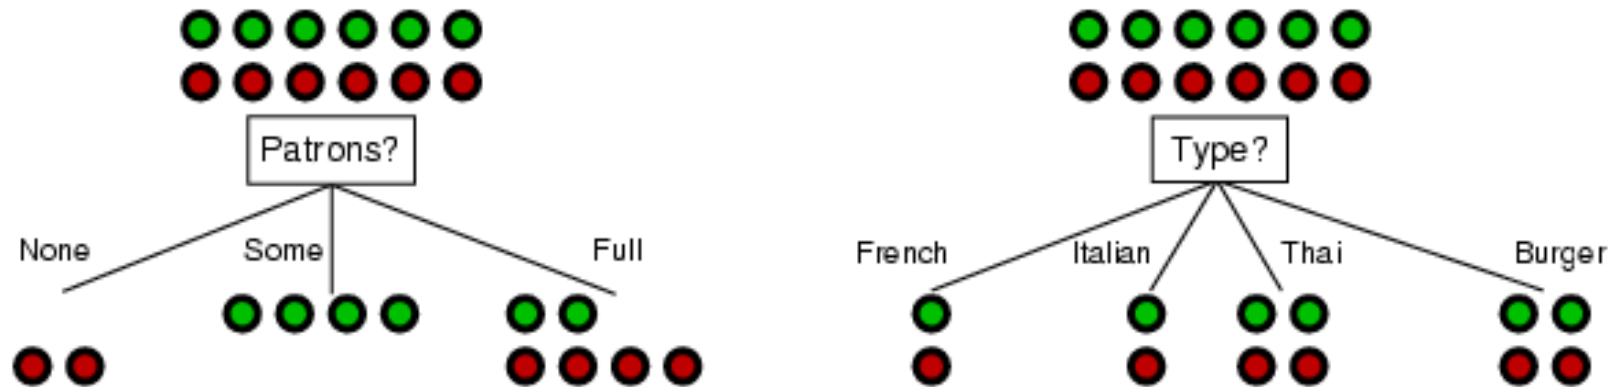
- no more input features
- all examples are classified the same
- too few examples to make an informative split

### 2. Which test to split on

- split gives smallest error.

# Choosing an attribute

Idea: good attribute splits examples into subsets that are (ideally) *all positive* or *all negative*

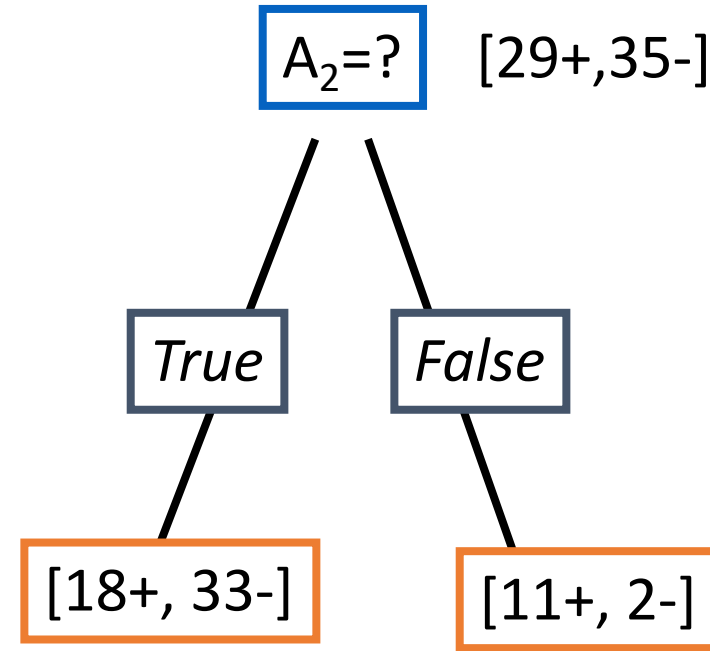
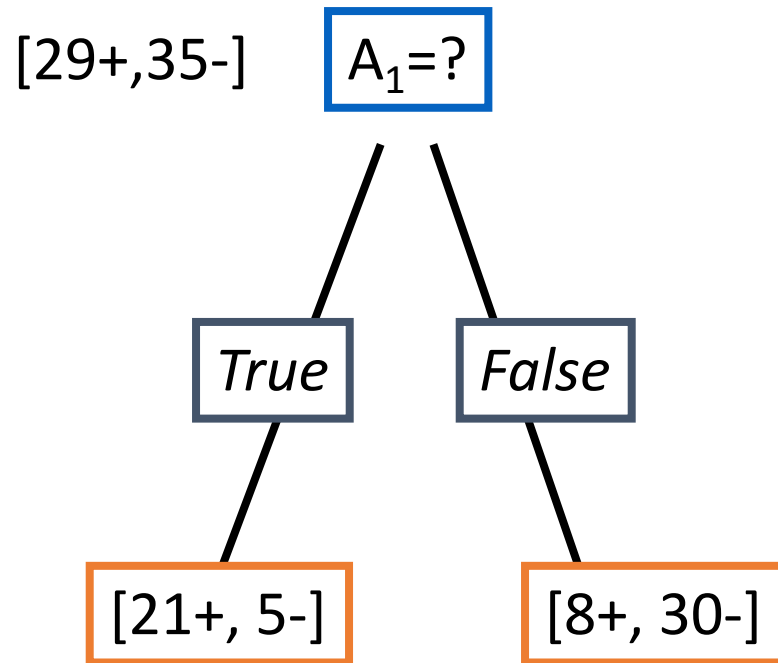


Which is better: *Patrons?* or *Type?*

Heuristic based on “**information gain**”.

- We want attributes that split the examples to sets that are relatively pure in one label; this way we are closer to a leaf node.

# Which Attribute is "best"?



# How to determine the Best Split?

- *Greedy approach:*
  - Nodes with *homogeneous class distribution* are preferred
- *Need a measure of node impurity:*

C0: 5
C1: 5

**Non-homogeneous,  
High degree of impurity**

C0: 9
C1: 1

**Homogeneous,  
Low degree of impurity**

**Information gain:** measures how well a given attribute separates the training examples according to their target classification

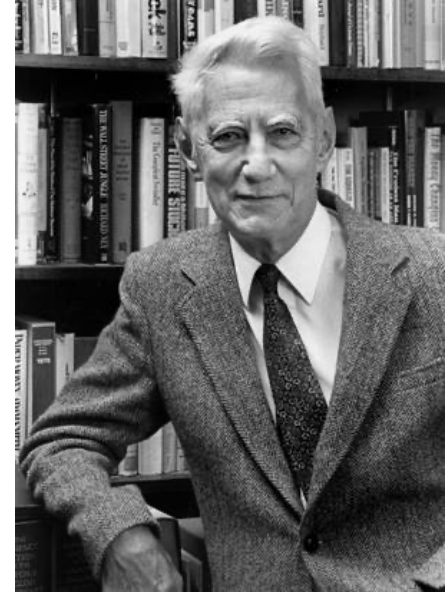
**Gini Index:** At each node measures, what is the error if you use the most prevalent label



# Background: Information theory

- Claude Shannon's seminal work: Mathematical Theory of Communication in 1948
- Information in a message (information entropy)
  - → minimum #bits needed to store/send using a **good** encoding
- If probability distribution  $P(p_1, p_2, \dots, p_n)$  for  $n$  messages, its information (or *entropy*) is:

$$H(P) = -(p_1 \log p_1 + p_2 \log p_2 + \dots + p_n \log p_n)$$



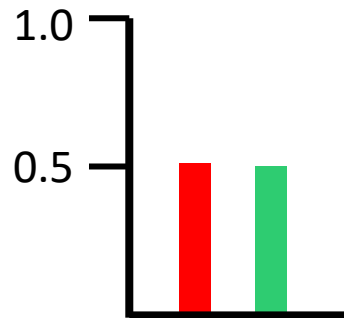
Claude Shannon

# Entropy of a Distribution

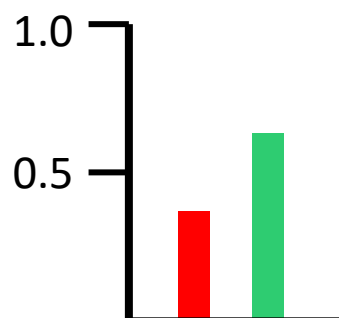
- Quantifies the amount of uncertainty associated with a specific probability distribution

$$H(X) = \sum_x p(X = x) \log_2 \frac{1}{p(X = x)}$$

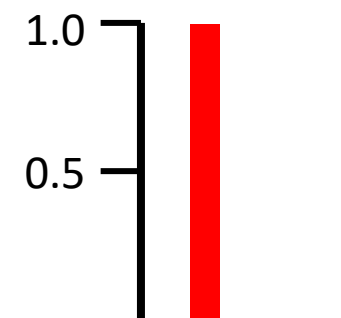
$$H(X) = - \sum_x p(X = x) \log_2 p(X = x)$$



$$\begin{aligned} H(P) &= -0.5 * (-1) + 0.5 \\ &\quad * (-1) \\ &= 1 \end{aligned}$$



$$\begin{aligned} H(P) &= -\left(\frac{2}{3} * \log\left(\frac{2}{3}\right) + \frac{1}{3} * \log\left(\frac{1}{3}\right)\right) \\ &= 0.92 \end{aligned}$$



$$\begin{aligned} H(P) &= -(1 * 1 + 0 * \log(0)) \\ &= 0 \end{aligned}$$

# Entropy (Measure I)

A measure for  
uncertainty  
purity  
information content

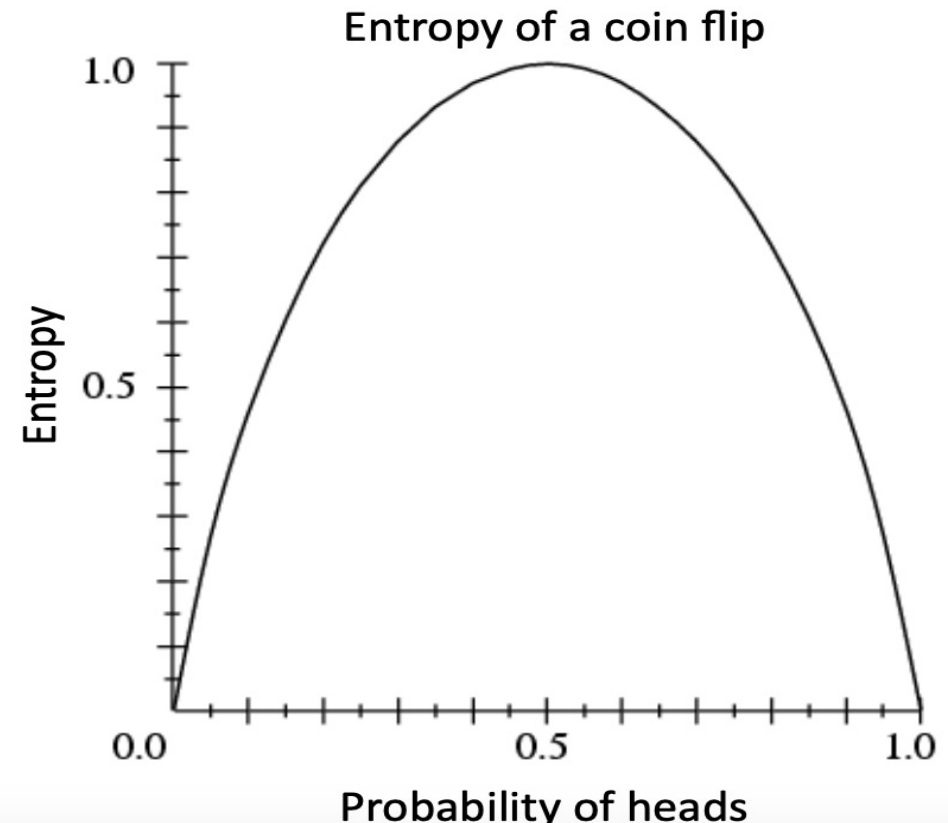
Entropy  $H(Y)$  of a random variable  $Y$

$$H(Y) = - \sum_{i=1}^k p(Y = y_i) \log_2 p(Y = y_i)$$

***More uncertainty, more entropy!***

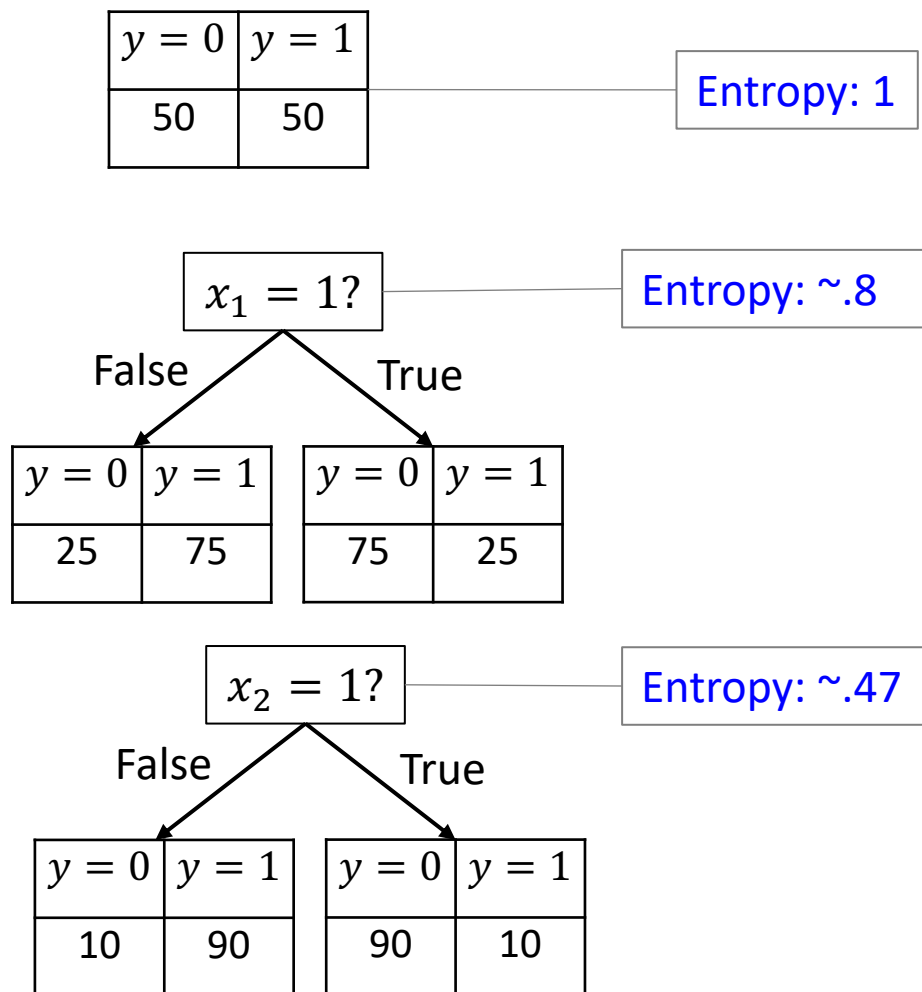
*Information Theory interpretation:*

$H(Y)$  is the expected number of bits needed to encode a randomly drawn value of  $Y$  (under most efficient code)



# Loss for Decision Trees

Should we stop splitting (in the first place)? Or continue?



$$Loss(S) = \frac{1}{n} \sum_{i=1}^n Entropy(Leaf(S_i))$$

Information Gain  
 $\sim .2$

Information Gain  
 $\sim .33$



**Information Gain** – reduction in Entropy (loss) from a change to the model

# Information Gain

Gain(S,A): expected reduction in entropy due to splitting S on attribute A

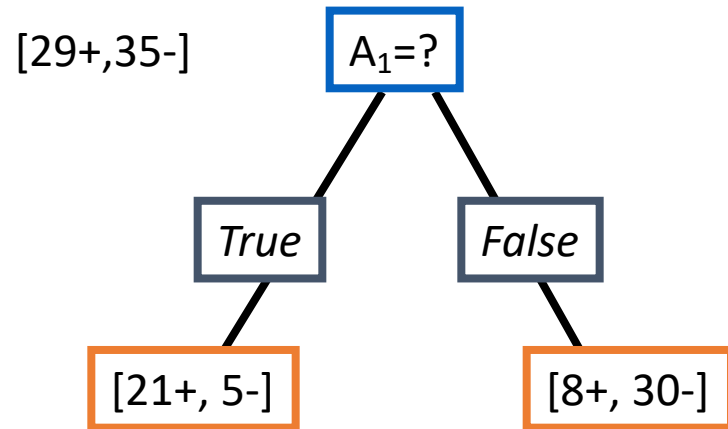
$$Gain(S, A) = Entropy(S) - \sum_{v \in values(A)} \frac{|S_v|}{|S|} \times Entropy(S_v)$$

$S_v$  is the subset of  $S$  for which attribute  $A$  has value  $v$ , and



$$\begin{aligned} Entropy([29+, 35-]) \\ &= -\frac{29}{64} \log_2 \frac{29}{64} - \frac{35}{64} \log_2 \frac{35}{64} \\ &= 0.99 \end{aligned}$$

# Information Gain Computation

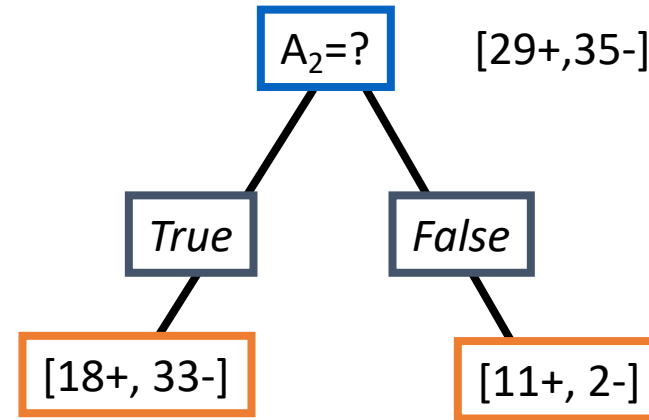


$$\text{Entropy}([21+, 5-]) = 0.71$$

$$\text{Entropy}([8+, 30-]) = 0.74$$

$$\text{Gain}(S, A_1) = \text{Entropy}(S)$$

$$\begin{aligned}
 & - \frac{26}{64} * \text{Entropy}([21+, 5-]) \\
 & - \frac{38}{64} * \text{Entropy}([8+, 30-]) \\
 & = 0.27
 \end{aligned}$$



$$\text{Entropy}([18+, 33-]) = 0.94$$

$$\text{Entropy}([8+, 30-]) = 0.62$$

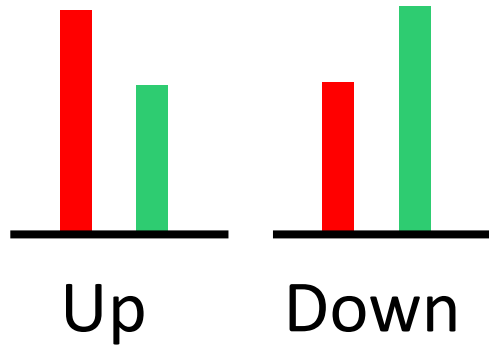
$$\text{Gain}(S, A_2)$$

$$\begin{aligned}
 & = \text{Entropy}(S) - \frac{51}{64} * \text{Entropy}([18+, 33-]) \\
 & - \frac{13}{64} * \text{Entropy}([11+, 2-]) \\
 & = 0.12
 \end{aligned}$$

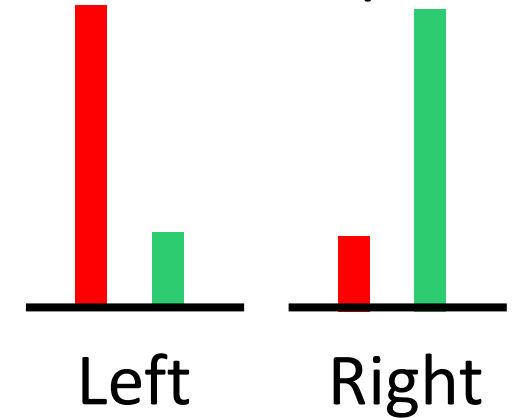
$$\begin{aligned}
 & \text{Entropy}([29+, 35-]) \\
 & = - \frac{29}{64} \log_2 \frac{29}{64} - \frac{35}{64} \log_2 \frac{35}{64} \\
 & = 0.99
 \end{aligned}$$

# An Illustration: DT with Real-Valued Features

“Best” (purest possible)  
Horizontal Split



“Best” (purest possible)  
Vertical Split



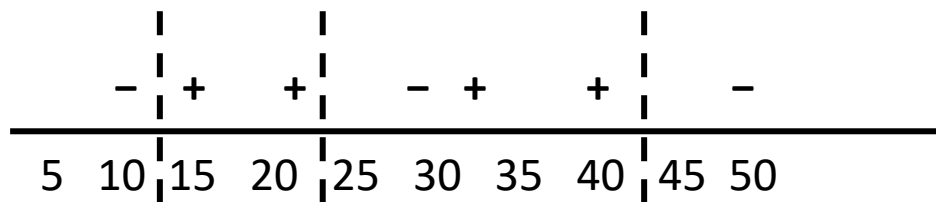
Between the best horizontal vs best vertical split, the vertical split is better (purer), hence we use this rule for the internal node

# DT with Real-Valued Features

Example:

- Length (L): 10 15 21 28 32 40 50
- Class:        -   +   +   -   +   +   -
- Check thresholds:

$L < 12.5$ ;  $L < 24.5$ ;  $L < 45$



How to find the **split with the highest gain** ?

For each continuous feature A:

- Sort examples according to the value of A
- For each ordered pair  $(x,y)$  with different labels
- Check the mid-point as a possible threshold.



# Gini Impurity: Decision Trees (Measure II)

- Gini impurity estimates the following
  - Choose an element randomly.
  - Label it using the distribution of labels on the set.
  - How often the element is incorrectly labeled?
- For a set of items with  $C$  classes, with relative frequency  $p_c$  for class  $c$ , the probability of choosing an item with label  $c$  is  $p_c$ , and the probability of misclassifying the item is

$$\sum_{k \neq c} p_k = 1 - p_c$$

- The Gini Index is computed by summing pairwise products of these probabilities for each class label:

$$GINI(p) = \sum_{c=1}^C p_c(1 - p_c) = 1 - \sum_{c=1}^C p_c^2$$

# Examples of Computing GINI

$$GINI(t) = 1 - \sum_j [p(j | t)]^2$$

C1	<b>0</b>
C2	<b>6</b>

$$P(C1) = 0/6 = 0 \quad P(C2) = 6/6 = 1$$

$$Gini = 1 - P(C1)^2 - P(C2)^2 = 1 - 0 - 1 = 0$$

C1	<b>1</b>
C2	<b>5</b>

$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$Gini = 1 - (1/6)^2 - (5/6)^2 = 0.278$$

C1	<b>2</b>
C2	<b>4</b>

$$P(C1) = 2/6 \quad P(C2) = 4/6$$

$$Gini = 1 - (2/6)^2 - (4/6)^2 = 0.444$$

# Measure of Impurity: Gini Index

Gini Index for a given node  $t$  :

$$GINI(t) = 1 - \sum_j [p(j | t)]^2$$

$p(j | t)$  is the relative frequency of class  $j$  at node  $t$

- **Maximum** ( $1 - 1/n_c$ ) when records are equally distributed among all classes, implying least interesting information [ $n_c$ : number of classes]
- **Minimum** (**0.0**) when all records belong to one class, implying most interesting information

C1	<b>0</b>
C2	<b>6</b>
<b>Gini=0.000</b>	

C1	<b>1</b>
C2	<b>5</b>
<b>Gini=0.278</b>	

C1	<b>2</b>
C2	<b>4</b>
<b>Gini=0.444</b>	

C1	<b>3</b>
C2	<b>3</b>
<b>Gini=0.500</b>	

# Splitting Based on GINI

Used in CART, SLIQ, SPRINT.

When a node  $p$  is split into  $k$  partitions (children), the quality of split is computed as,

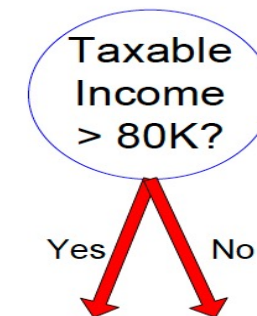
$$GINI_{split} = \sum_{i=1}^k \frac{n_i}{n} GINI(i)$$

where,  $n_i$  = number of records at child  $i$ ,  
 $n$  = number of records at node  $p$ .

# Continuous Attributes: Computing Gini Index

- Use binary decisions ( $A < v$  and  $A \geq v$ )
- Many choice.
- Each value  $v$  has a count matrix associated with it
  - Class counts with  $A \geq v$
  - Class counts with  $A < v$
- Method to choose best  $v$ 
  - For each  $v$ , scan the dataset to get the count matrix. Compute Gini Index.
  - Choose the one with min. Gini Index.
  - **Inefficient. Repetition of work!**

Tid	Refund	Marital Status	Taxable Income	Cheat
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes





# Continuous Attributes: Computing Gini Index...

- For efficient computation: for each attribute,
  - Sort the attribute on values
  - Linearly scan these values, each time updating the count matrix and computing gini index
  - Choose the split position that has the least gini index

		Cheat	No		No		No		Yes		Yes		Yes		No		No		No		No		
		Taxable Income																					
Sorted Values Split Positions		60		70		75		85		90		95		100		120		125		220			
		55		65		72		80		87		92		97		110		122		172		230	
		<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>	<=	>
	Yes	0	3	0	3	0	3	0	3	1	2	2	1	3	0	3	0	3	0	3	0	3	0
	No	0	7	1	6	2	5	3	4	3	4	3	4	3	4	4	3	5	2	6	1	7	0
	Gini	0.420		0.400		0.375		0.343		0.417		0.400		<u>0.300</u>		0.343		0.375		0.400		0.420	

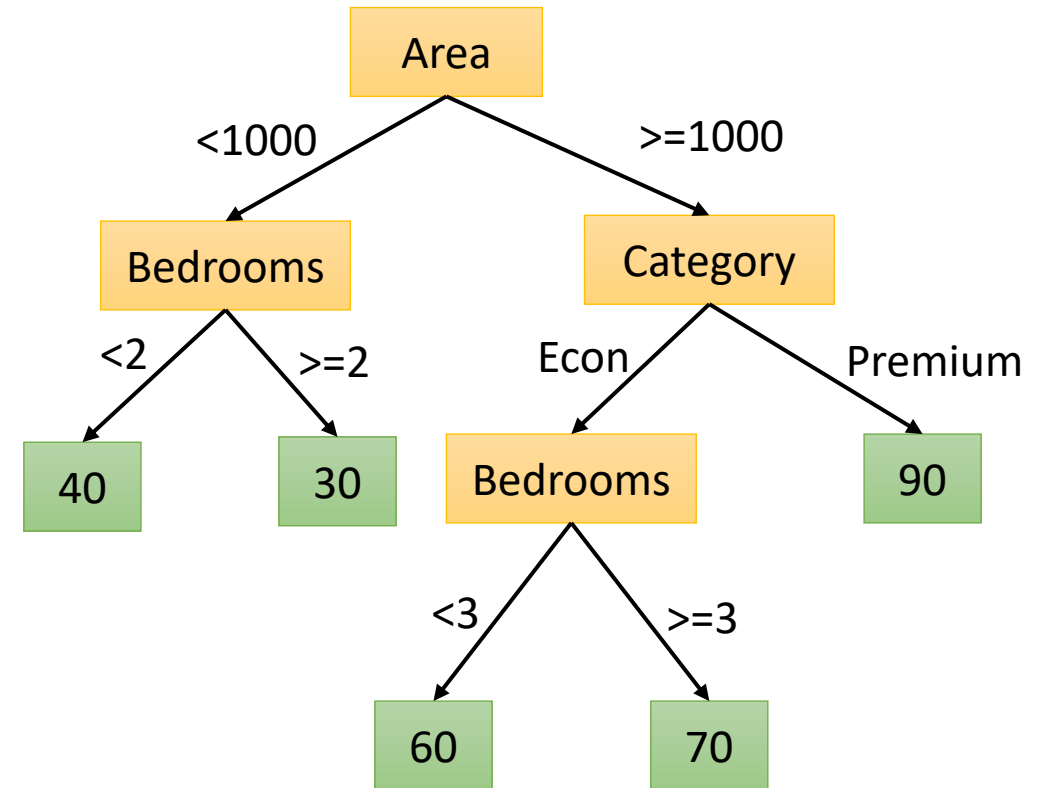
# Continuous Attributes: Variance (Measure III)

Another Example of Splitting Criteria:

## Variance Reduction

If a node is entirely homogeneous, then the variance is zero.

1. For each split, individually calculate the variance of each child node
2. Calculate the variance of each split as the weighted average variance of child nodes
3. Select the split with the lowest variance



# Practical Issues

- Missing Values
- Attributes with different costs
  - Change information gain so that low cost attribute are preferred
- Dealing with features with different # of values



# Experimental Machine Learning

- Machine Learning is an Experimental Field
- Basics:
  - Split your data into two (or three) sets:
    - Training data (often 70-90%)
    - Test data (often 10-20%)
    - Development data (10-20%)
- You need to report performance on test data, but you are not allowed to look at it.
  - You are allowed to look at the development data (and use it to tweak parameters)